

# Code for AI 3 Lab Work 1

## Loading the Dataset

```
import pandas as pd
```

```
df = pd.read_csv("./dataset/Wine_Test_01.csv")  
df.head(10)
```

```
X = df.drop("quality", axis=1)  
y = df["quality"]
```

## Task 1

```
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

```
from sklearn.svm import SVC  
from sklearn.metrics import accuracy_score  
from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import GridSearchCV, train_test_split
```

```
# a) Display a histogram for each attribute
```

```
def plot_all_histograms(X):  
    X.hist(bins=20, figsize=(15, 10))  
    plt.suptitle("Histograms of All Attributes", y=1.02)  
    plt.tight_layout()  
    plt.show()
```

```
# b) Plot histograms of each attribute separated by target Y
```

```
def plot_histograms_by_class(X):  
    for column in X.columns:  
        plt.figure(figsize=(6, 4))  
        sns.histplot(data=df, x=column, hue="quality", kde=True, element="step", stat="density",  
common_norm=False)  
        plt.title(f"Distribution of {column} by Wine Quality")  
        plt.tight_layout()  
        plt.show()
```

```

# c) Perform 10 runs of SVM classification with grid search
def run_svm_classification(df):
    X = df.drop("quality", axis=1)
    y = df["quality"]

    param_grid = {
        'C': [0.1, 1, 10],
        'kernel': ['linear', 'rbf'],
        'gamma': ['scale', 0.01, 0.001]
    }

    accuracies = []
    best_params_list = []

    for i in range(10):
        print(f"\nRun {i+1}")
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=i)

        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_test_scaled = scaler.transform(X_test)

        grid = GridSearchCV(SVC(), param_grid, cv=5, scoring='accuracy')
        grid.fit(X_train_scaled, y_train)

        best_model = grid.best_estimator_
        y_pred = best_model.predict(X_test_scaled)
        acc = accuracy_score(y_test, y_pred)

        print("Best Params:", grid.best_params_)
        print("Accuracy:", round(acc, 4))

        best_params_list.append(grid.best_params_)
        accuracies.append(acc)

    avg_acc = np.mean(accuracies)
    std_acc = np.std(accuracies)

    print("\n==== Summary =====")
    print(f"Average Accuracy: {avg_acc:.4f}")
    print(f"Standard Deviation: {std_acc:.4f}")

    plot_all_histograms(X)

```

```
plot_histograms_by_class(X)
run_svm_classification(df)
```

## Task 2

# a) Choose an attribute with overlap – we use 'volatile acidity'

```
def plot_density(attribute='volatile acidity'):
    plt.figure(figsize=(6, 4))
    sns.histplot(data=df, x=attribute, hue="quality", kde=True, element="step", stat="density",
common_norm=False)
    plt.title(f"Original Distribution: {attribute}")
    plt.tight_layout()
    plt.show()
```

# b) Delete samples to increase class separation

```
def reduce_overlap(attribute='volatile acidity', lower=0.3, upper=0.6):
    df_reduced = df[(df[attribute] < lower) | (df[attribute] > upper)]
    return df_reduced
```

```
plot_density()
df_reduced = reduce_overlap(attribute='volatile acidity', lower=0.3, upper=0.6)
print(f"Reduced data size: {len(df_reduced)} samples")
# c) Rerun classification like Task 1c
run_svm_classification(df_reduced)
```

## Task 3

```
from sklearn.decomposition import PCA
```

```
def plot_cumulative_variance(df):
    # 1. Standardize the dataset
    X_train, _, _, _ = train_test_split(X, y, test_size=0.2, random_state=42)
    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)

    # 2. Use the PCA class that does PCA on the dataset
    pca = PCA()
    _ = pca.fit_transform(X_train_scaled)

    # 3. Find and plot the cumulative variance
    cumulative_variance = pca.explained_variance_ratio_.cumsum()
```

```

number_of_components = len(cumulative_variance)

plt.figure(figsize=(8, 5))
plt.plot(range(1, number_of_components+1), cumulative_variance)
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Variance')
plt.grid(True)
plt.show()

plot_cumulative_variance(df)

def run_svm_with_pca(df, n_components=8):
    X = df.drop("quality", axis=1)
    y = df["quality"]

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    # Apply PCA
    pca = PCA(n_components=n_components)
    X_train_pca = pca.fit_transform(X_train_scaled)
    X_test_pca = pca.transform(X_test_scaled)

    # SVM with Grid Search
    param_grid = {
        'C': [0.1, 1, 10],
        'kernel': ['linear', 'rbf'],
        'gamma': ['scale', 'auto']
    }

    svm = SVC()
    grid_search = GridSearchCV(svm, param_grid, cv=5)
    grid_search.fit(X_train_pca, y_train)

    best_params = grid_search.best_params_
    y_pred = grid_search.predict(X_test_pca)
    accuracy = accuracy_score(y_test, y_pred)

    return accuracy, best_params

accuracy, best_params = run_svm_with_pca(df, n_components=8)
print("Best Hyperparameters:", best_params)

```

```
print(f"Test Accuracy with PCA + SVM: {accuracy:.4f}")
```

## Task 4

```
from sklearn.feature_selection import RFE
```

```
def run_svm_with_rfe(df, n_features=8):
```

```
    X = df.drop("quality", axis=1)
```

```
    y = df["quality"]
```

```
    # 1. Standardize the dataset
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

```
    scaler = StandardScaler()
```

```
    X_train_scaled = scaler.fit_transform(X_train)
```

```
    X_test_scaled = scaler.transform(X_test)
```

```
    base_estimator = SVC(kernel='linear', C=1000, gamma=0.1)
```

```
    selector = RFE(base_estimator, n_features_to_select=n_features)
```

```
    selector.fit(X_train_scaled, y_train)
```

```
    selected_mask = selector.support_
```

```
    selected_features = X.columns[selected_mask]
```

```
    print("Selected Features:", selected_features.tolist())
```

```
    # Reduce feature sets
```

```
    X_train_reduced = X_train[selected_features]
```

```
    X_test_reduced = X_test[selected_features]
```

```
    # Grid search on reduced features
```

```
    param_grid = {
```

```
        'C': [0.1, 1, 10],
```

```
        'kernel': ['linear', 'rbf'],
```

```
        'gamma': ['scale', 'auto']
```

```
    }
```

```
    grid = GridSearchCV(SVC(), param_grid, cv=5)
```

```
    grid.fit(X_train_reduced, y_train)
```

```
    best_params = grid.best_params_
```

```
    y_pred = grid.predict(X_test_reduced)
```

```
    accuracy = accuracy_score(y_test, y_pred)
```

```
    return selected_features.tolist(), best_params, accuracy
```

```
selected_features, best_params, accuracy = run_svm_with_rfe(df, n_features=8)
print("Best Hyperparameters:", best_params)
print(f"Test Accuracy with RFE + SVM: {accuracy:.4f}")
```